

Building Crawlers Using Built-in Services in Scrapy



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Scrapy has multiple features which are useful when crawling websites at scale

Logging events to console and to file

Using the telnet utility to debug crawlers

Working with broad crawlers to concurrently crawl multiple sites

Auto throttling crawls so as to not fall afoul of website policies

Demo


Logging to console and file

Demo

**Sending email notifications based on
crawl results**


Broad Crawls

Connected Graph with No Cycle




Such a graph is called a **tree**

Connected Graph with No Cycle



Trees are great for depicting
hierarchical relationships

Tree



Trees are great for depicting
hierarchical relationships

Two Ways of Traversing Graphs

Breadth-first

All nodes at same distance from origin visited together

Depth-first

All nodes in certain direction from origin visited together

Crawling web pages is quite similar to traversing a tree where the initial web page is the root node

Two Ways of Traversing Graphs


Breadth-first

All nodes at same distance from origin visited together

Depth-first


All nodes in certain direction from origin visited together

“Breadth-first” Tree Traversal




Nodes are visited level-by-level

“Breadth-first” Tree Traversal




Visited H

“Breadth-first” Tree Traversal




Visited H - B

“Breadth-first” Tree Traversal




Visited H - B - F

“Breadth-first” Tree Traversal




Visited H - B - F - A

“Breadth-first” Tree Traversal




Visited H - B - F - A - G

“Breadth-first” Tree Traversal




Visited H - B - F - A - G - E

“Breadth-first” Tree Traversal



Visited H - B - F - A - G - E - C

“Breadth-first” Tree Traversal



Visited H - B - F - A - G - E - C - D

Two Ways of Traversing Graphs


Breadth-first

All nodes at same distance from
origin visited together


Depth-first

All nodes in certain direction from
origin visited together

“Depth-first” Tree Traversal




“Depth-first” Tree Traversal




Visited H

“Depth-first” Tree Traversal




Visited H - B

“Depth-first” Tree Traversal



Visited H - B - A


“Depth-first” Tree Traversal



In breadth-first, would have visited F before A


Visited H - B - A

“Depth-first” Tree Traversal




Visited H - B - A

“Depth-first” Tree Traversal




Visited H - B - A - F

“Depth-first” Tree Traversal




Visited H - B - A - F - G

“Depth-first” Tree Traversal




Visited H - B - A - F - G - E

“Depth-first” Tree Traversal



Visited H - B - A - F - G - E - C

“Depth-first” Tree Traversal



Visited H - B - A - F - G - E - C

Two Ways of Traversing Graphs

Breadth-first

All nodes at same distance from origin visited together

Depth-first

All nodes in certain direction from origin visited together

Two Ways of Using Scrapy

Broad Crawls

Crawl a forest of domains, stopping only when resources exhausted; used in search engines

Depth-first

All nodes in certain direction from origin visited together

Two Ways of Using Scrapy

Broad Crawls

Crawl a forest of domains, stopping only when resources exhausted; used in search engines

Focused Crawls

Crawl a specific site, usually with a single Scrapy Spider

Most Scrapy defaults are for focused crawls, which are similar logically to depth-first traversal

Two Types of Crawls

Focused Crawls

Usually crawl a single, specific domain

Crawl to completion

Logic is usually quite complex - handled inside single Spider

Parallelism in scraping a single domain restricted by politeness

Broad Crawls

Usually used to crawl an unbounded set of domains

Limit crawl by number of pages or time

Logic is usually simple - data processed in a separate stage

Parallelism in scraping multiple domains is not limited, hence very fast

Broad Crawls are very fast
because the full parallelism of
Scrapy comes into play

Broad Crawl Best Practices

Increase global concurrency

Increase max thread pool size

Set up your own DNS

Reduce log level

Disable cookies

Disable retries

Reduce download timeout

Disable redirects

Enable crawl of Ajax “crawlable pages”

Broad Crawl Best Practices

Increase global concurrency

Increase max thread pool size

Set up your own DNS

Reduce log level

Disable cookies

Disable retries

Reduce download timeout

Disable redirects

Enable crawl of Ajax “crawlable pages”

Broad Crawl Best Practices

Increase global concurrency

Increase max thread pool size

Set up your own DNS

Reduce log level

Disable cookies

Disable retries

Reduce download timeout

Disable redirects

Enable crawl of Ajax “crawlable pages”

Broad Crawl Best Practices

Increase global concurrency

Increase max thread pool size

Set up your own DNS

Reduce log level

Disable cookies

Disable retries

Reduce download timeout

Disable redirects

Enable crawl of Ajax “crawlable pages”

Broad Crawl Best Practices

Increase global concurrency

Increase max thread pool size

Set up your own DNS

Reduce log level

Disable cookies

Disable retries

Reduce download timeout

Disable redirects

Enable crawl of Ajax “crawlable pages”

Demo

Building broad crawlers using Scrapy

Demo

Debugging crawlers using telnet

Demo

Auto throttling broad crawlers

Summary

Scrapy has multiple features which are useful when crawling websites at scale

Logging events to console and to file

Using the telnet utility to debug crawlers

Working with broad crawlers to concurrently crawl multiple sites

Auto throttling crawls so as to not fall afoul of website policies